

Daniel Andres Jimenez Riveros

Cloud AI Engineer | AWS, Large Language Models & Inference Optimization

Vancouver, BC, Canada | da.jimenez166@gmail.com | linkedin.com/in/daniel-jimenez-78500117a | github.com/daanjiri

Professional Summary

Cloud AI / ML Engineer with 5+ years of experience deploying, adapting, and optimizing Large Language Models (LLMs) on AWS (SageMaker, Bedrock, EC2, Fargate, Lambda). Hands-on with domain-specific fine-tuning (LoRA, PEFT, SFT, RLHF/DPO), prompt engineering, quantization, and inference acceleration for latency and cost tuning. Ships production RAG and multi-modal systems in PyTorch and Hugging Face, backed by MLOps for model versioning, monitoring, and ML CI/CD.

Work Experience

Apexon, Vancouver, BC, Canada

Senior GenAI Engineer (Cloud AI, LLM Deployment on AWS)

June 2025 - Present

- **Projected +\$0.8B net new revenue from prioritized ISV verticals in the first launch year** by designing, deploying, and serving an agentic LLM application for a payment platform (McKinsey client) on AWS with Python and LangGraph, using tool and function calling, multi-agent orchestration, and guardrails behind containerized FastAPI services.
- **Achieved 96% accuracy in a production RAG system** by architecting retrieval pipelines on AWS (Bedrock, OpenSearch, S3, Lambda) with Anthropic Claude and OpenAI models and advanced prompt and context engineering, evaluated with DeepEval and Ragas and monitored in Langfuse for quality, latency, and token cost.
- **Cut inference latency and cost per request** by right-sizing and routing models across Bedrock and third-party LLM APIs, compressing prompts and context, caching, and streaming; as **technical lead**, established model and prompt versioning, automated evaluation gates, ML CI/CD, and infrastructure as code.

Arkangel AI, Remote

Senior Machine Learning Engineer (LLM Fine-Tuning & Model Serving)

June 2024 - June 2025

- **45,000+ users served by an LLM agent for clinical text extraction** from electronic health records, built with Hugging Face Transformers, RAG, and LangGraph and deployed on AWS Fargate, ECR, and S3 with Docker, FastAPI, and autoscaling for high-availability inference.
- **Raised task accuracy through domain-specific model adaptation** using supervised fine-tuning (SFT), LoRA/PEFT adapters, and RLHF/DPO alignment on medical corpora in PyTorch and Hugging Face, benchmarked on held-out evaluation sets before promoting each model version.
- **Generated up to \$50k monthly recurring revenue** with an AutoML and deep learning platform on AWS SageMaker, PyTorch, Lambda, S3, and CloudWatch: training jobs, managed endpoints, model registry and versioning, automated retraining, and drift monitoring.
- **Reduced processing time and cost by 70%** with multi-modal document AI using vision-language models (VLMs), OCR, and computer vision (OpenCV, PyTorch).

Amazon, Remote

Software Development Engineer

July 2022 - August 2023

- **Analyzed 500,000+ customer service transactions for 9,000+ Amazon employees** with Python (Pandas, NumPy) text analytics and automated data pipelines, and **empowered 300+ e-commerce leaders** with scalable backend services and RESTful APIs in TypeScript and NodeJS, delivered through code review, testing, and CI/CD.
-

Selected Projects

Thorcast AI, On-Premise LLM Platform for Financial Services

2025 - Present

- **Generating \$5k+ monthly revenue** from banks with an application that builds credit and operational risk ML/DL models in minutes instead of months, automating ingestion, feature engineering, training, and evaluation with an AI agent in Python, FastAPI, LangGraph, and PyTorch, and **serving quantized open-weight LLMs on-premise with Ollama** under strict data-privacy constraints.

Alzheimer's Detection from Speech, UBC Master's Thesis

March 2024 - August 2024

- **Achieved 90% recall** by fine-tuning transformer models and CNNs in PyTorch and TensorFlow for multi-modal analysis of speech and imaging, with hyperparameter tuning and transfer learning.
-

Education

Master of Data Science in Computational Linguistics, University of British Columbia, Vancouver, Canada (2023-2024)

Master of Science in Information Technology, Universidad de los Andes, Bogota, Colombia (2018-2020)

Bachelor of Engineering in Industrial Engineering and Bachelor of Arts, Universidad de los Andes (2014-2018)

Awards: COLFUTURO Scholarship (2023) | Data Science for All, DS4A (2018)

Skills

AWS & LLM Deployment: Amazon SageMaker (Training Jobs, Real-Time & Asynchronous Endpoints, Model Registry), Amazon Bedrock, EC2, Lambda, Fargate/ECS, ECR, S3, OpenSearch, CloudWatch, Auto Scaling, Docker, Kubernetes, FastAPI, Ollama, Azure

Fine-Tuning & Inference Optimization: Domain-Specific Fine-Tuning, LoRA, PEFT, Supervised Fine-Tuning (SFT), RLHF, DPO Alignment, Prompt & Context Engineering, Quantization, Inference Acceleration, Latency Tuning, Cost Optimization, Model Routing & Right-Sizing, Batching, Caching, Model Serving

MLOps: Model Versioning & Model Registry, Model Monitoring & Drift Detection, ML CI/CD Pipelines (GitHub Actions), Automated Retraining, MLflow, Weights & Biases, Langfuse Observability, Terraform/IaC, A/B Testing

ML & GenAI: PyTorch, Hugging Face Transformers, PEFT/TRL, TensorFlow, Scikit-learn, LangChain, LangGraph, RAG Architectures, Vector Databases (OpenSearch, FAISS, Pinecone, ChromaDB), Multi-Modal Models (VLMs, OCR), NLP, LLM Evaluation (DeepEval, Ragas), Guardrails

Programming & Engineering: Python (Expert), SQL, TypeScript, JavaScript, C++, Java, C#/.NET, Git & GitHub, RESTful APIs, Microservices, Testing, Code Review